Aviation Accident Prediction Using Deep Learning

Group Members:

Marc George - 43
Merlin Sarah Jiju - 45
Timon K John - 67

INTRODUCTION

- Despite technological advances, aviation accidents still occur due to complex interactions among human, environmental, and mechanical factors.
- Early identification of potential risks can help prevent accidents and enhance overall flight safety.
- Deep learning models can process vast aviation data—such as flight parameters, weather, and maintenance logs—to uncover hidden patterns and predict anomalies.
- Hence the aim is to develop an intelligent predictive system that minimizes aviation risks through accurate and data-driven accident forecasting.

OBJECTIVE

1. Evaluate the effectiveness of deep neural networks in analyzing complex aviation flight data, including sensor telemetry and operational metadata.
2. Identify critical environmental, technical, and human factors influencing accident likelihood, leveraging domain-driven feature extraction and representation learning.
3. Develop a conceptual predictive framework demonstrating how such models can assist pilots, engineers, and safety managers in making timely, informed safety decisions.
4. Assess the practical benefits and limitations of deploying deep learning models for real-world aviation safety risk mitigation and proactive intervention.

DATASETS USED

- Synthetic Dataset - The data is a synthetic, programmatically generated dataset .
- The data contents emulates flight telemetry, weather, phase, and binary incident outcomes (fatal vs non-fatal).
- All features and labels are constructed to simulate realistic diversity and balance, not directly sourced from real-world aviation records.

DATA CLEANING AND INTEGRATION

- No conventional cleaning or deduplication is required due to synthetic data generation.
- Metadata and sequence telemetry are constructed for each sample and combined without missing values or inconsistencies.

EXPLORATORY DATA ANALYSIS

- EDA includes inspecting class distribution, variable shapes, and simulated feature distributions for sanity checking.
- Explored correlations among weather, flight phase, and aircraft characteristics to identify patterns influencing accident likelihood.
- All class counts and distributions are controlled by design, confirming intended balance.

LABEL GENERATION

- Outcome labels are directly assigned during synthetic sample creation: fatal = 1, non-fatal = 0.
- Train and validation splits are stratified by label.

FEATURE EXTRACTION & DATA PREPARATION

- **Metadata Features**: Weather, flight phase, and maintenance data were transformed into model-ready formats using *scikit-learn* encoders (e.g., one-hot and label encoding).
- **Scaling and Normalization:** Continuous variables (e.g., altitude, airspeed, temperature) were standardized using *MinMaxScaler* and *StandardScaler* to ensure consistent numeric ranges, improving convergence during training.
- **Sequential Telemetry Construction:** Time-series flight telemetry (sensor readings over time) was reshaped into **sequence tensors**, capturing the temporal dependencies between successive readings. This mimics the real-world flight progression from takeoff to landing.

MODEL ARCHITECTURE - AVIATION TRANSFORMER

1. **Input Layer**- The model accepts **two parallel inputs**:

Sequential Telemetry Data – multivariate time-series sensor readings (e.g., altitude, velocity, pitch, temperature).
Metadata Inputs – contextual features like weather, flight phase, and maintenance status.

2. **Positional Encoding -** Since transformer models do not inherently understand sequence order, **positional encodings** are added to telemetry tensors.

These encodings help the model recognize the temporal sequence of flight readings (e.g., takeoff → cruise → landing).

3. **Transformer Encoder Block -**

Multi-Head Self-Attention: Allows the model to focus on important time steps or sensor variables that most influence risk.

Feed-Forward Neural Network: Applies nonlinear transformations to capture complex relationships.

Layer Normalization & Residual Connections: Maintain gradient stability and smooth information flow across layers.

MODEL ARCHITECTURE - AVIATION TRANSFORMER

4. **Sequence Pooling:**

The encoder outputs for each time step are aggregated (via **mean or attention pooling**) into a single vector summarizing the flight's entire temporal behavior.

5. **Metadata Head (MLP Block):**

A **Multi-Layer Perceptron (MLP)** processes non-sequential contextual data (e.g., aircraft type, weather, maintenance) which consists of one or more dense layers with ReLU activations to learn nonlinear patterns from categorical and continuous metadata.

6. **Fusion Layer:**

The pooled sequence embedding from the Transformer and the processed metadata vector from the MLP are **concatenated** to create a unified feature representation of the flight event.

7. **Classification Layer:**

A **fully connected dense layer** with a **sigmoid activation** outputs the probability of the flight resulting in a *fatal* or *non-fatal*

MODEL TRAINING AND OPTIMIZATION

- Loss: Focal Loss (handles balanced or imbalanced data by default).
- Optimizer: Ranger (AdamW-like) with OneCycleLR scheduling.
- Regularization: Dropout and weight decay.
- Epochs: 100
- Early stopping based on validation loss [after 48 epochs]
- Hyperparameter tuning for transformer depth, attention heads, batch size.

MODEL EVALUATION

- Metrics: Accuracy, Precision, Recall, F1-Score, ROC-AUC, PR-AUC are computed.
- A confusion matrix and classification report are produced post-validation.

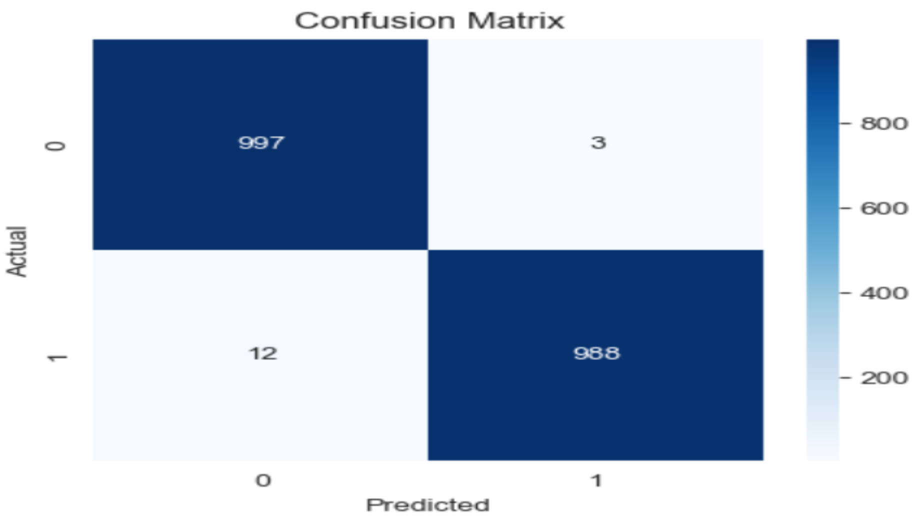

PERFORMANCE VISUALIZATION

- Accuracy and loss tracked across epochs for training and validation sets.

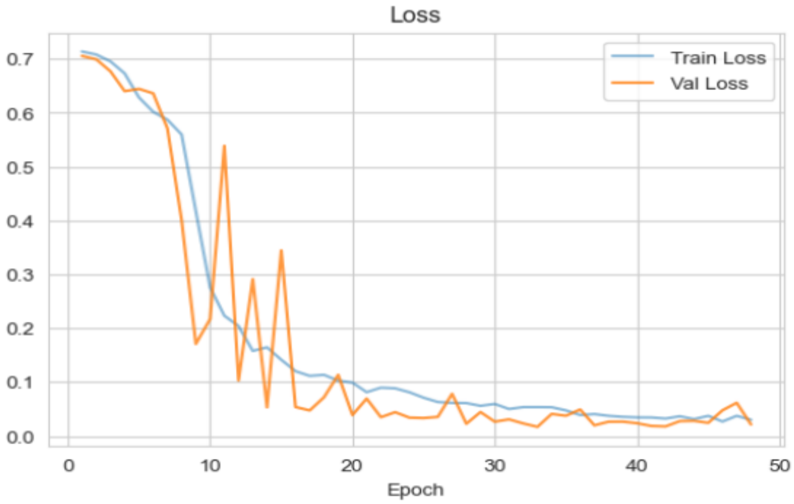

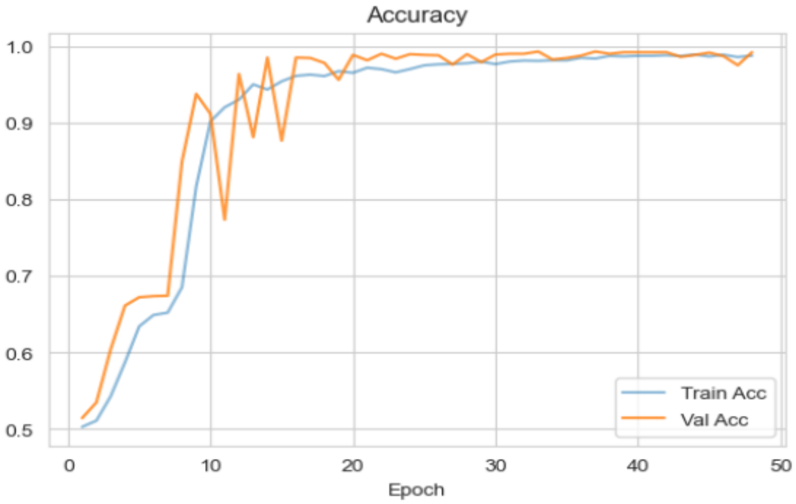

PERFORMANCE VISUALIZATION

- High AUC(0.99) confirms strong discriminative capability
- PR curve demonstrates balanced precision-recall tradeoff

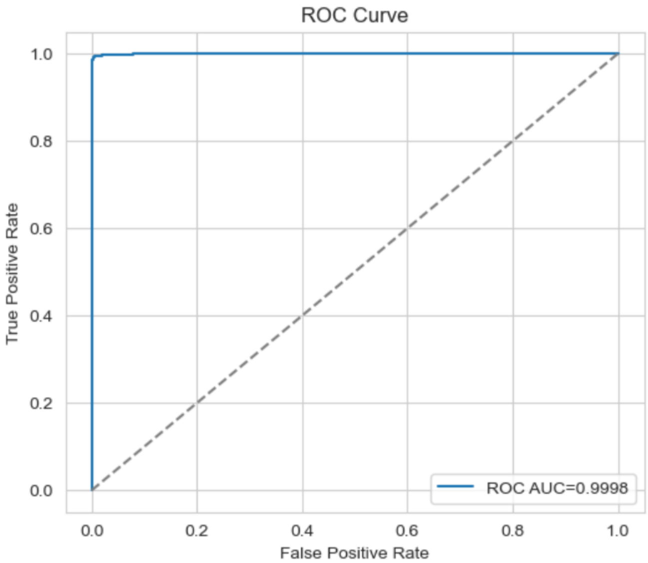

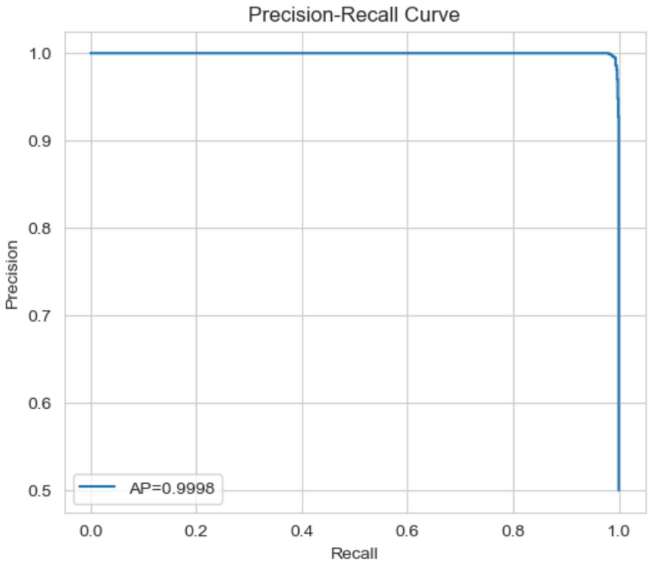

RISK SCORING & CATEGORIZATION

- Predicted probability outputs categorized into:
  - Low Risk (0.00 to 0.33)
  - Medium Risk (> 0.33 to 0.66)
  - High Risk (> 0.66 to 1.00)

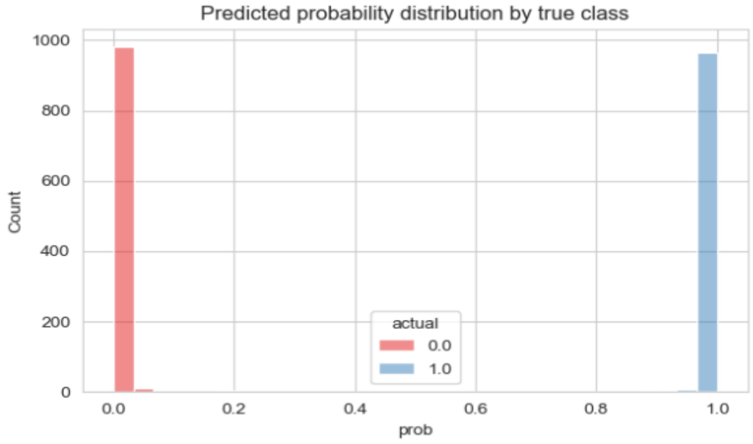

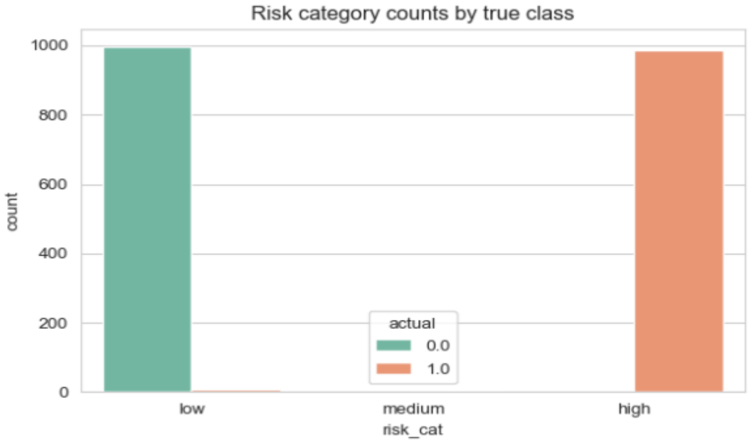

RESUTS SUMMARY

- Test Accuracy: 0.9925 (99.25%)
- ROC – AUC: 0.99

	prob	pred	actual	risk_cat	meta_cat__weather_clear	meta_cat__weather_fog	meta_cat__weather_heavy_turbuler
0	0.090334	0	0.0	low	1.0	0.0	
1	1.000000	1	1.0	high	1.0	0.0	
2	0.000058	0	0.0	low	0.0	0.0	
3	0.999956	1	1.0	high	0.0	0.0	
4	0.000020	0	0.0	low	0.0	0.0	
5	0.000192	0	0.0	low	1.0	0.0	
6	0.000743	0	0.0	low	0.0	0.0	
7	1.000000	1	1.0	high	1.0	0.0	
8	0.000092	0	0.0	low	0.0	0.0	
9	1.000000	1	1.0	high	0.0	0.0	
10	0.001200	0	0.0	low	0.0	0.0	

```
Result summary:
  Samples: 2000
  Accuracy: 0.9925  Precision: 0.9970  Recall: 0.9880  F1: 0.9925
  ROC AUC: 0.9998  Average Precision (AP): 0.9998
```

CONCLUSION

- The Transformer-based model effectively captured complex temporal and contextual relationships within aviation telemetry and metadata.
- The integration of metadata (weather, flight phase, maintenance) with sequential flight data improved model interpretability and overall predictive power.
- The model achieved high recall for fatal incidents and demonstrated strong, consistent performance under stratified validation, confirming reliable learning behavior.
- The experiment validates deep learning's potential for structured sequential event prediction and proactive aviation safety management.
- Results are based on synthetic data, indicating proof of concept; real-world deployment would require validation on authentic aviation datasets for operational applicability.